



Real time Analytics with Apache Kafka and Spark

October 2014 Meetup

Organized by **Big Data Hyderabad.**

<http://www.meetup.com/Big-Data-Hyderabad/>

Rahul Jain
 @rahuldausa

About Me

- Big data/Search Consultant
- 8+ years of learning experience.
- Worked (got a chance) on High volume distributed applications.
- Still a learner (and beginner)

Quick Questionnaire

How many people know/work on Scala ?

How many people know/work on Apache Kafka?

How many people know/heard/are using Apache Spark ?

What we are going to learn/see today ?

- Apache Zookeeper (Overview)
- Apache Kafka – Hands-on
- Apache Spark – Hands-on
- Spark Streaming (Explore)




Apache Zookeeper™

What is Zookeeper



- An Open source, High Performance coordination service for distributed applications
- Centralized service for
 - Configuration Management
 - Locks and Synchronization for providing coordination between distributed systems
 - Naming service (Registry)
 - Group Membership
- Features
 - hierarchical namespace
 - provides watcher on a znode
 - allows to form a cluster of nodes
- Supports a large volume of request for data retrieval and update
- <http://zookeeper.apache.org/>



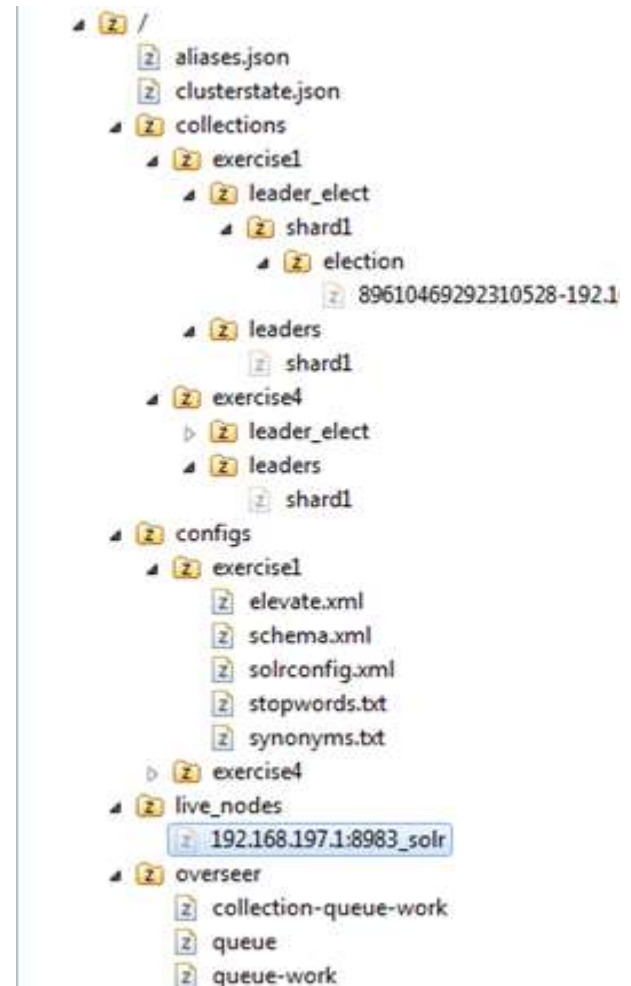
Source : <http://zookeeper.apache.org>

Zookeeper Use cases

- Configuration Management
 - Cluster member nodes Bootstrapping configuration from a central source
- Distributed Cluster Management
 - Node Join/Leave
 - Node Status in real time
- Naming Service – e.g. DNS
- Distributed Synchronization – locks, barriers
- Leader election
- Centralized and Highly reliable Registry


Zookeeper Data Model

- Hierarchical Namespace
- Each node is called “znode”
- Each znode has data(stores data in byte[] array) and can have children
- znode
 - Maintains “Stat” structure with version of data changes , ACL changes and timestamp
 - Version number increases with each changes




Let's recall basic concepts of
Messaging System

Point to Point Messaging (Queue)



Publish-Subscribe Messaging (Topic)

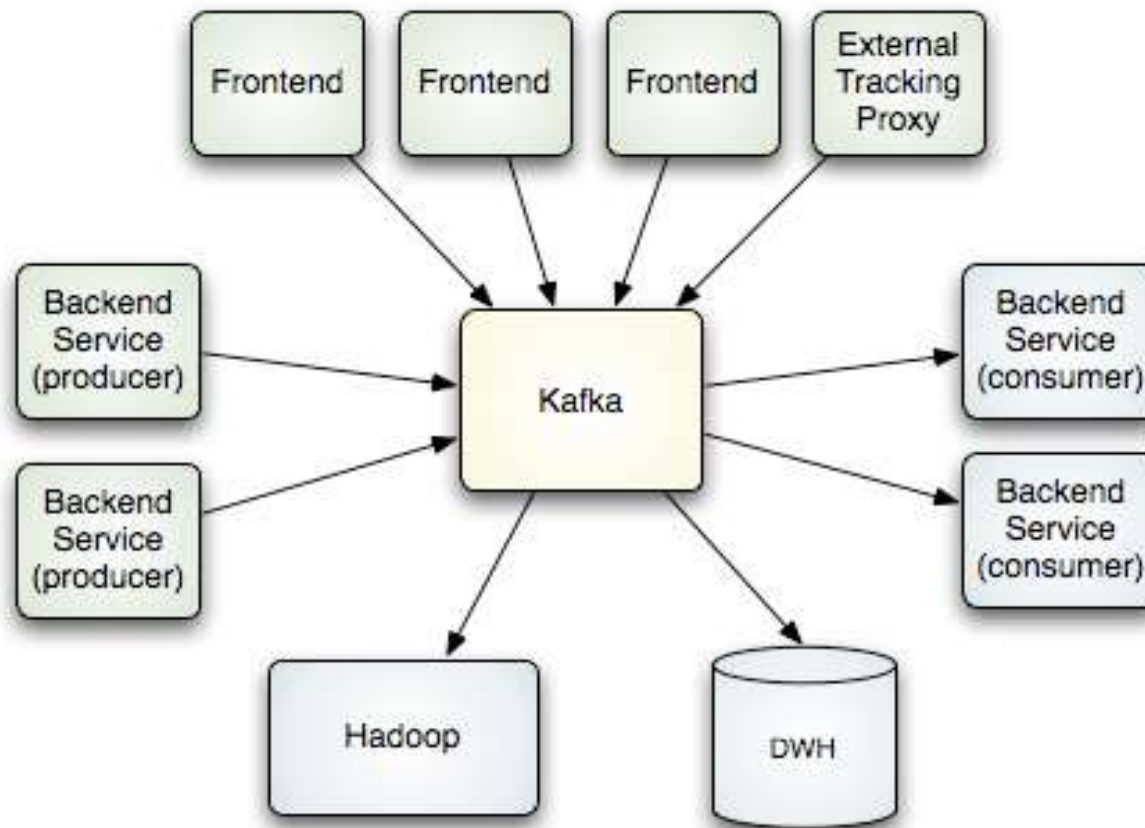


Apache Kafka

Overview

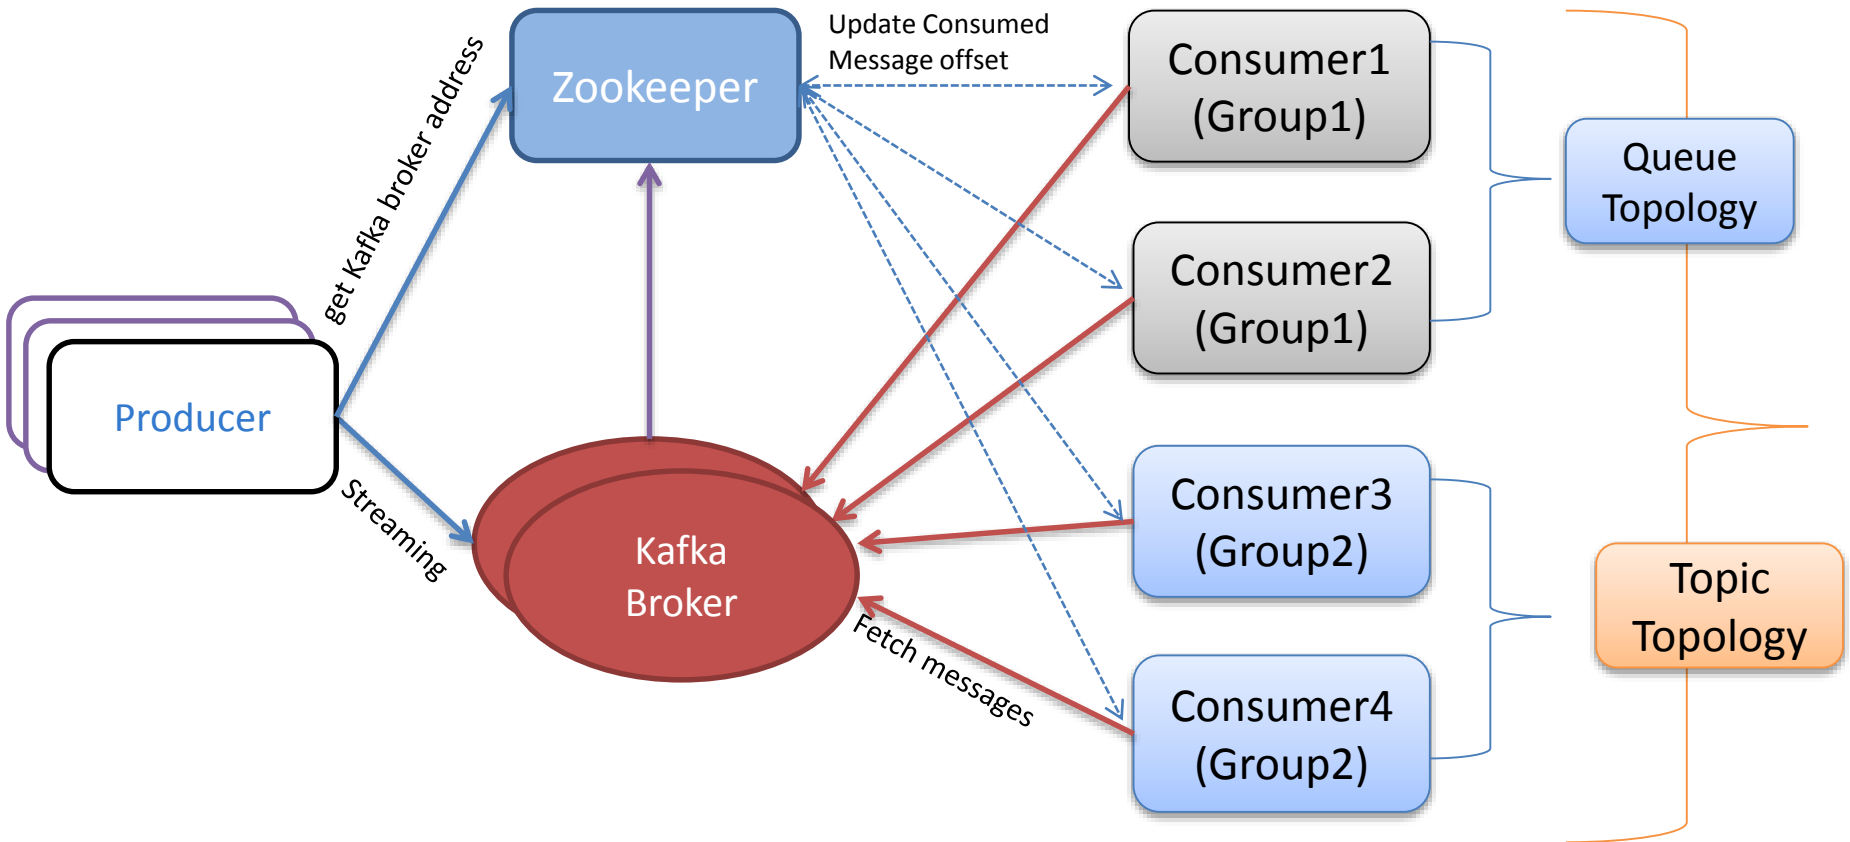
- An apache project initially developed at LinkedIn
- Distributed publish-subscribe messaging system
- Designed for processing of real time activity stream data e.g. logs, metrics collections
- Written in Scala
- Does not follow JMS Standards, neither uses JMS APIs
- Features
 - Persistent messaging
 - High-throughput
 - Supports both queue and topic semantics
 - Uses Zookeeper for forming a cluster of nodes (producer/consumer/broker)and many more...
- <http://kafka.apache.org/>

How it works




Real time transfer

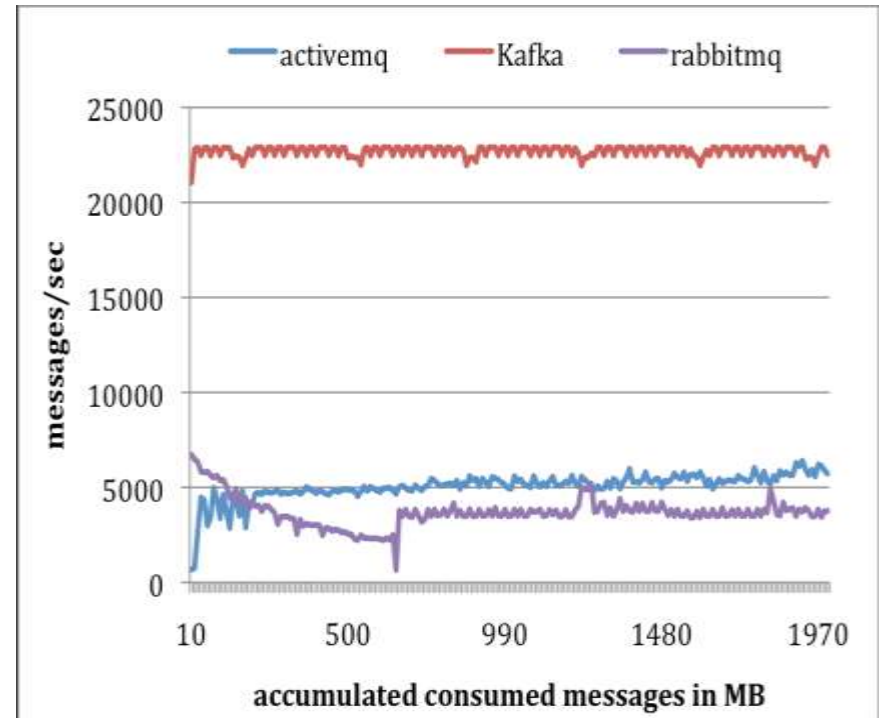
Broker does not **Push** messages to Consumer, Consumer **Polls** messages from Broker.



Performance Numbers



Producer Performance



Consumer Performance

About Apache Spark




- Initially started at UC Berkeley in 2009
- Fast and general purpose cluster computing system
- 10x (on disk) - 100x (In-Memory) faster
- Most popular for running *Iterative Machine Learning Algorithms*.
- Provides high level APIs in
 - Java
 - Scala
 - Python
- Integration with Hadoop and its eco-system and can **read existing data**.
- <http://spark.apache.org/>


So Why Spark ?

- Most of Machine Learning Algorithms are iterative because each iteration can improve the results
- With Disk based approach each iteration's output is written to disk making it slow

Hadoop execution flow




Spark execution flow



<http://www.wiziq.com/blog/hype-around-apache-spark/>


Spark Stack

- Spark SQL
 - For SQL and unstructured data processing
- MLib
 - Machine Learning Algorithms
- GraphX
 - Graph Processing
- Spark Streaming
 - stream processing of live data streams



<http://spark.apache.org>

Execution Flow



Terminology

- **Application Jar**
 - User Program and its dependencies except Hadoop & Spark Jars bundled into a Jar file
- **Driver Program**
 - The process to start the execution (main() function)
- **Cluster Manager**
 - An external service to manage resources on the cluster (standalone manager, YARN, Apache Mesos)
- **Deploy Mode**
 - **cluster** : Driver inside the cluster
 - **client** : Driver outside of Cluster

Terminology (contd.)

- **Worker Node** : Node that run the application program in cluster
- **Executor**
 - Process launched on a worker node, that runs the Tasks
 - Keep data in memory or disk storage
- **Task** : A unit of work that will be sent to executor
- **Job**
 - Consists multiple tasks
 - Created based on a Action
- **Stage** : Each Job is divided into smaller set of tasks called Stages that is sequential and depend on each other
- **SparkContext** :
 - represents the connection to a Spark cluster, and can be used to create RDDs, accumulators and broadcast variables on that cluster.

Resilient Distributed Dataset (RDD)

- Resilient Distributed Dataset (RDD) is a basic Abstraction in Spark
- Immutable, Partitioned collection of elements that can be operated in parallel
- Basic Operations
 - map
 - filter
 - persist
- Multiple Implementation
 - [PairRDDFunctions](#) : RDD of Key-Value Pairs, groupByKey, Join
 - [DoubleRDDFunctions](#) : Operation related to double values
 - [SequenceFileRDDFunctions](#) : Operation related to SequenceFiles
- RDD main characteristics:
 - A list of partitions
 - A function for computing each split
 - A list of dependencies on other RDDs
 - Optionally, a Partitioner for key-value RDDs (e.g. to say that the RDD is hash-partitioned)
 - Optionally, a list of preferred locations to compute each split on (e.g. block locations for an HDFS file)
- Custom RDD can be also implemented (by overriding functions)

Cluster Deployment

- Standalone Deploy Mode
 - simplest way to deploy Spark on a private cluster
- Amazon EC2
 - EC2 scripts are available
 - Very quick launching a new cluster
- Apache Mesos
- Hadoop YARN

Monitoring

The screenshot shows a web browser window with the title "Spark shell - Spark Stages". The address bar shows "localhost:4040/stages/". The page has a navigation bar with the Spark logo and tabs for "Stages", "Storage", "Environment", and "Executors". The "Stages" tab is active. The main content area is titled "Spark Stages" and displays the following information:

- Total Duration: 45 s
- Scheduling Mode: FIFO
- Active Stages: 0
- Completed Stages: 0
- Failed Stages: 0

Below this information, there are three sections, each with a table header:

Active Stages (0)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
----------	-------------	-----------	----------	------------------------	--------------	---------------

Completed Stages (0)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
----------	-------------	-----------	----------	------------------------	--------------	---------------

Failed Stages (0)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write	Failure Reason
----------	-------------	-----------	----------	------------------------	--------------	---------------	----------------

Monitoring – Stages

Spark Pi - Details for Stage x

localhost4040/stages/stage/?id=0

Spark

StagesStorageEnvironmentExecutors

Spark Pi application UI

Details for Stage 0

Total task time across all tasks: 41 s

Summary Metrics for 73 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Result serialization time	0 ms	0 ms	0 ms	0 ms	21 ms
Duration	0.2 s	0.4 s	0.5 s	0.6 s	1.0 s
Time spent fetching task results	0 ms	0 ms	0 ms	0 ms	0 ms
Scheduler delay	5 ms	9 ms	14 ms	29 ms	1 s

Aggregated Metrics by Executor

Executor ID	Address	Task Time	Total Tasks	Failed Tasks	Succeeded Tasks	Shuffle Read	Shuffle Write	Shuffle Spill (Memory)	Shuffle Spill (Disk)
localhost	CANNOT FIND ADDRESS	47 s	70	0	70	0.0 B	0.0 B	0.0 B	0.0 B

Tasks

Task Index	Task ID	Status	Locality Level	Executor	Launch Time	Duration	GC Time	Result Ser Time	Errors
0	0	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.5 s			
1	1	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.3 s		21 ms	
2	2	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.8 s	98 ms	1 ms	
3	3	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.5 s			
4	4	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.7 s	98 ms		
5	5	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.7 s	98 ms		
6	6	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.7 s	98 ms	1 ms	
7	7	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:05	0.8 s	98 ms		
8	8	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:07	0.6 s	98 ms	1 ms	
9	9	SUCCESS	PROCESS_LOCAL	localhost	2014/09/13 02:22:07	0.3 s			



Let's start getting hands dirty

Kafka Installation

- Download
 - <http://kafka.apache.org/downloads.html>
- Untar it
 - > `tar -xzf kafka_<version>.tgz`
 - > `cd kafka_<version>`

Start Servers

- Start the Zookeeper server
 - > `bin/zookeeper-server-start.sh config/zookeeper.properties`

Pre-requisite: Zookeeper should be up and running.

- Now Start the Kafka Server
 - > `bin/kafka-server-start.sh config/server.properties`

Create/List Topics

- Create a topic

```
> bin/kafka-topics.sh --create --zookeeper localhost:2181 --  
replication-factor 1 --partitions 1 --topic test
```

- List down all topics

```
> bin/kafka-topics.sh --list --zookeeper localhost:2181
```

Output: test

Producer

- Send some Messages

```
> bin/kafka-console-producer.sh --broker-list  
localhost:9092 --topic test
```

Now type on console:

This is a message

This is another message

Consumer

- Receive some Messages

```
> bin/kafka-console-consumer.sh --zookeeper  
localhost:2181 --topic test --from-beginning
```

This is a message

This is another message

Multi-Broker Cluster

- Copy configs

- > cp config/server.properties config/server-1.properties
 - > cp config/server.properties config/server-2.properties

- Changes in the config files.

config/server-1.properties:

broker.id=1

port=9093

log.dir=/tmp/kafka-logs-1

config/server-2.properties:

broker.id=2

port=9094

log.dir=/tmp/kafka-logs-2

Start with New Nodes

- Start other Nodes with new configs

```
> bin/kafka-server-start.sh config/server-1.properties &  
> bin/kafka-server-start.sh config/server-2.properties &
```

- Create a new topic with replication factor as 3

```
> bin/kafka-topics.sh --create --zookeeper  
localhost:2181 --replication-factor 3 --partitions  
1 --topic my-replicated-topic
```

- List down the all topics

```
> bin/kafka-topics.sh --describe --zookeeper  
localhost:2181 --topic my-replicated-topic
```

Topic:my-replicated-topic PartitionCount:1 ReplicationFactor:3 Configs:

Topic: my-replicated-topic Partition: 0 Leader: 1 Replicas: 1,2,0 Isr: 1,2,0



Let's move to
Apache Spark

Spark Shell

```
./bin/spark-shell --master local[2]
```

The `--master` option specifies the [master URL for a distributed cluster](#), or `local` to run locally with one thread, or `local[N]` to run locally with N threads. You should start by using `local` for testing.

```
./bin/run-example SparkPi 10
```

This will run 10 iterations to calculate the value of Pi

Basic operations...

```
scala> val textFile = sc.textFile("README.md")
textFile: spark.RDD[String] = spark.MappedRDD@2ee9b6e3
```

```
scala> textFile.count() // Number of items in this RDD
res0: Long = 126
```

```
scala> textFile.first() // First item in this RDD
res1: String = # Apache Spark
```

```
scala> val linesWithSpark = textFile.filter(line =>
line.contains("Spark"))
linesWithSpark: spark.RDD[String] = spark.FilteredRDD@7dd4af09
```

Simplier - Single liner:

```
scala> textFile.filter(line => line.contains("Spark")).count()
// How many lines contain "Spark"?
res3: Long = 15
```

Map - Reduce

```
scala> textFile.map(line => line.split(" ").size).reduce((a, b)
=> if (a > b) a else b)
res4: Long = 15
```

```
scala> import java.lang.Math
scala> textFile.map(line => line.split(" ").size).reduce((a, b)
=> Math.max(a, b))
res5: Int = 15
```

```
scala> val wordCounts = textFile.flatMap(line => line.split("
")).map(word => (word, 1)).reduceByKey((a, b) => a + b)
wordCounts: spark.RDD[(String, Int)] =
spark.ShuffledAggregatedRDD@71f027b8
wordCounts.collect()
```

With Caching...

```
scala> linesWithSpark.cache()  
res7: spark.RDD[String] = spark.FilteredRDD@17e51082
```

```
scala> linesWithSpark.count()  
res8: Long = 15
```

```
scala> linesWithSpark.count()  
res9: Long = 15
```

With HDFS...

```
val lines = spark.textFile("hdfs://...")  
val errors = lines.filter(line => line.startsWith("ERROR"))  
println(Total errors: + errors.count())
```


Standalone (Scala)

```
/* SimpleApp.scala */
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._
import org.apache.spark.SparkConf

object SimpleApp {

  def main(args: Array[String]) {
    val logFile = "YOUR_SPARK_HOME/README.md" // Should be some file on your
system
    val conf = new SparkConf().setAppName("Simple Application")
.setMaster("local")
    val sc = new SparkContext(conf)
    val logData = sc.textFile(logFile, 2).cache()
    val numAs = logData.filter(line => line.contains("a")).count()
    val numBs = logData.filter(line => line.contains("b")).count()
    println("Lines with a: %s, Lines with b: %s".format(numAs, numBs))
  }
}
```

Standalone (Java)

```
/* SimpleApp.java */
import org.apache.spark.api.java.*;
import org.apache.spark.SparkConf;
import org.apache.spark.api.java.function.Function;

public class SimpleApp {

    public static void main(String[] args) {

        String logFile = "YOUR_SPARK_HOME/README.md"; // Should be some file on your system
        SparkConf conf = new SparkConf().setAppName("Simple Application").setMaster("local");
        JavaSparkContext sc = new JavaSparkContext(conf);
        JavaRDD<String> logData = sc.textFile(logFile).cache();

        long numAs = logData.filter(new Function<String, Boolean>() {
            public Boolean call(String s) { return s.contains("a"); }
        }).count();

        long numBs = logData.filter(new Function<String, Boolean>() {
            public Boolean call(String s) { return s.contains("b"); }
        }).count();

        System.out.println("Lines with a: " + numAs + ", lines with b: " + numBs);
    }
}
```

Job Submission

```
$SPARK_HOME/bin/spark-submit \  
  --class "SimpleApp" \  
  --master local[4] \  
  target/scala-2.10/simple-project_2.10-1.0.jar
```

Configuration

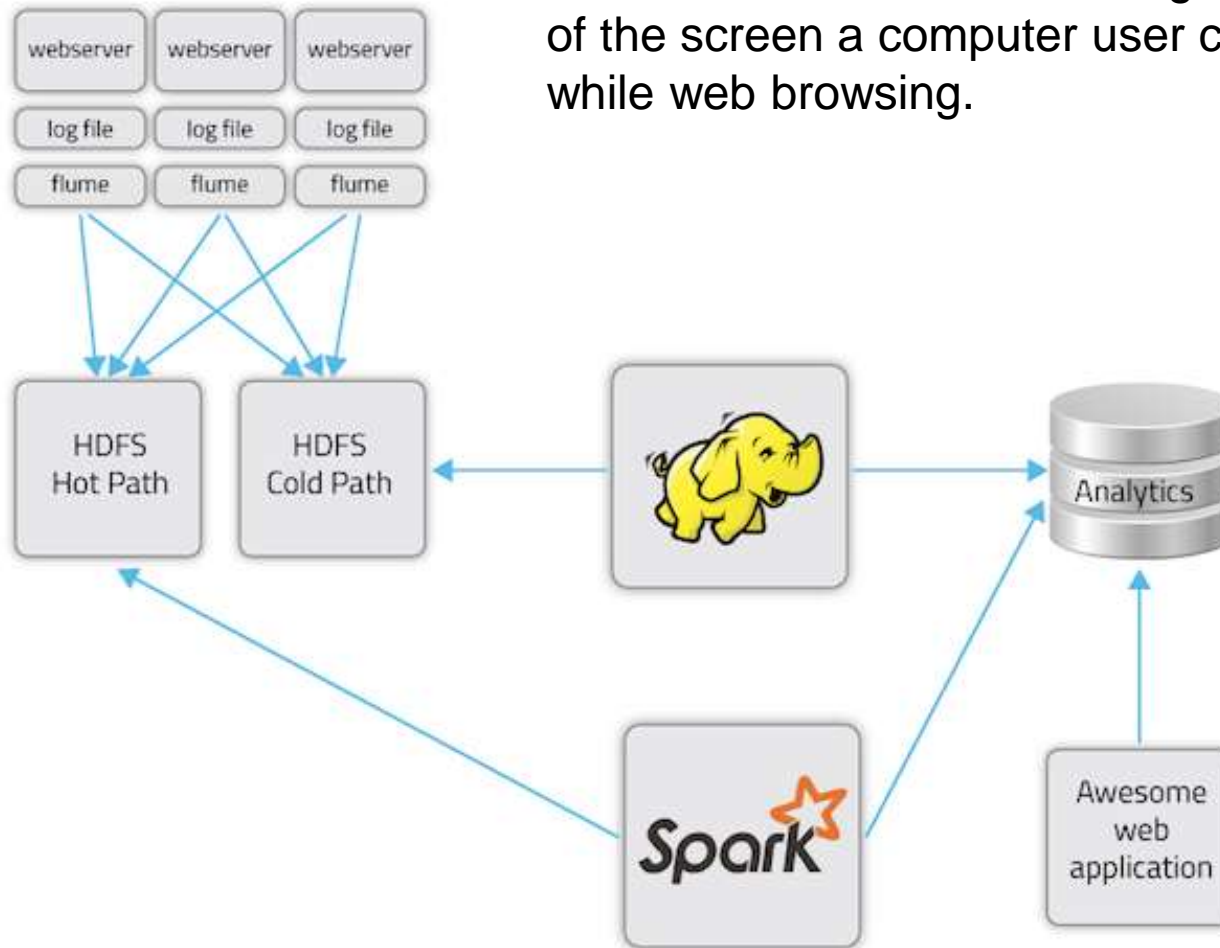
```
val conf = new SparkConf()  
    .setMaster("local")  
    .setAppName("CountingSheep")  
    .set("spark.executor.memory", "1g")  
  
val sc = new SparkContext(conf)
```

Let's discuss

Real-time Clickstream Analytics

Analytics High Level Flow

A **clickstream** is the recording of the parts of the screen a computer user clicks on while web browsing.



Spark Streaming

Makes it easy to build scalable fault-tolerant streaming applications.

- Ease of Use
- Fault Tolerance
- Combine streaming with batch and interactive queries.

Support Multiple Input/Outputs




Streaming Flow



Spark Streaming Terminology

- Spark Context
 - An object with configuration parameters
- Discretized Steams (DStreams)
 - Stream from input sources or generated after transformation
 - Continuous series of RDDs
- Input DStreams
 - Stream of raw data from input sources



Spark Context

```
SparkConf sparkConf = new SparkConf().setAppName("PageViewsCount");  
sparkConf.setMaster("local[2]");//running locally with 2 threads  
  
// Create the context with a 5 second batch size  
JavaStreamingContext jssc = new JavaStreamingContext(sparkConf, new Duration(5000));
```

Reading from Kafka Stream

```
int numThreads = Integer.parseInt(args[3]);  
Map<String, Integer> topicMap = new HashMap<String, Integer>();  
String[] topics = args[2].split(",");  
for (String topic: topics) {  
    topicMap.put(topic, numThreads);  
}  
JavaPairReceiverInputDStream<String, String> messages =  
    KafkaUtils.createStream(jssc, <zkQuorum>, <group>, topicMap);
```

Processing of Stream

```
/**
 * Incoming:
 * // (192.168.2.82,1412977327392,www.example.com,192.168.2.82)
 * // 192.168.2.82: key : tuple2._1()
 * // 1412977327392,www.example.com,192.168.2.82: value: tuple2._2()
 */
// Method signature: messages.map(InputArgs, OutputArgs)
JavaDStream<Tuple2<String, String>> events = messages.map(new Function<Tuple2<String, String>, Tuple2<String, String>>() {
    @Override
    public Tuple2<String, String> call(Tuple2<String, String> tuple2) {
        String[] parts = tuple2._2().split(",");
        return new Tuple2<>(parts[2], parts[1]);
    }
});
```

```
JavaPairDStream<String, Long> pageCountsByUrl = events.map(new Function<Tuple2<String,String>, String>(){
    @Override
    public String call(Tuple2<String, String> pageView) throws Exception {
        return pageView._2();
    }
}).countByValue();

System.out.println("Page counts by url:");

pageCountsByUrl.print();
```

Thanks!

@rahuldausa on twitter and slideshare
<http://www.linkedin.com/in/rahuldausa>

Join us @ For Solr, Lucene, Elasticsearch, Machine Learning, IR
<http://www.meetup.com/Hyderabad-Apache-Solr-Lucene-Group/>
<http://www.meetup.com/DataAnalyticsGroup/>

Join us @ For Hadoop, Spark, Cascading, Scala, NoSQL, Crawlers and all cutting edge technologies.
<http://www.meetup.com/Big-Data-Hyderabad/>